

import package and load data

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
df = pd.read_csv('/content/MachineLearningLab/diabetes.csv')
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6		0.627	50
1	1	85	66	29	0	26.6		0.351	31
2	8	183	64	0	0	23.3		0.672	32
3	1	89	66	23	94	28.1		0.167	21
4	0	137	40	35	168	43.1		2.288	33

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
# Replace zero values in key features with column median
key_features = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']

for col in key_features:
    median = df.loc[df[col] != 0, col].median()
    df.loc[df[col] == 0, col] = median

# Replace the first row's glucose value with the maximum glucose
df.loc[0, 'Glucose'] = df['Glucose'].max()

# For records with the lowest age, replace glucose values with the minimum
min_age = df['Age'].min()
df.loc[df['Age'] == min_age, 'Glucose'] = df['Glucose'].min()

df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	199	72	35	125	33.6		0.627	50
1	1	85	66	29	125	26.6		0.351	31
2	8	183	64	29	125	23.3		0.672	32
3	1	44	66	23	94	28.1		0.167	21
4	0	137	40	35	168	43.1		2.288	33

Next steps: [Generate code with df](#) [New interactive sheet](#)

```

print(df.info())
print(df.describe())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Pregnancies      768 non-null    int64  
 1   Glucose          768 non-null    int64  
 2   BloodPressure    768 non-null    int64  
 3   SkinThickness    768 non-null    int64  
 4   Insulin          768 non-null    int64  
 5   BMI              768 non-null    float64 
 6   DiabetesPedigreeFunction 768 non-null    float64 
 7   Age              768 non-null    int64  
 8   Outcome          768 non-null    int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
None
Pregnancies      Glucose      BloodPressure  SkinThickness  Insulin  \
count    768.000000  768.000000  768.000000  768.000000  768.000000
mean      3.845052  116.294271   72.386719   29.108073  140.671875
std       3.369578  36.797403   12.096642   8.791221  86.383060
min       0.000000  44.000000   24.000000   7.000000  14.000000
25%      1.000000  95.000000   64.000000   25.000000 121.500000
50%      3.000000  115.000000   72.000000   29.000000 125.000000
75%      6.000000  140.000000   80.000000   32.000000 127.250000
max      17.000000 199.000000  122.000000  99.000000 846.000000

          BMI  DiabetesPedigreeFunction  Age  Outcome
count  768.000000                  768.000000  768.000000
mean   32.455208                   0.471876  33.240885  0.348958
std    6.875177                   0.331329  11.760232  0.476951
min   18.200000                   0.078000  21.000000  0.000000
25%  27.500000                   0.243750  24.000000  0.000000
50%  32.300000                   0.372500  29.000000  0.000000
75%  36.600000                   0.626250  41.000000  1.000000
max   67.100000                  2.420000  81.000000  1.000000

```

```

X = df.drop('Outcome', axis=1)
y = df['Outcome']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)

model = LinearRegression()
model.fit(X_train, y_train)

y_pred_cont = model.predict(X_test)
y_pred = np.round(y_pred_cont).astype(int)
y_pred = np.clip(y_pred, 0, 1)

```

```

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, zero_division=0)
rec = recall_score(y_test, y_pred, zero_division=0)
f1 = f1_score(y_test, y_pred, zero_division=0)
cm = confusion_matrix(y_test, y_pred)

```

```
print("Accuracy:", acc)
print("Precision:", prec)
print("Recall:", rec)
print("F1-score:", f1)
print("Confusion Matrix:\n", cm)
```


```
Accuracy: 0.7597402597402597
Precision: 0.68
Recall: 0.6181818181818182
F1-score: 0.6476190476190476
Confusion Matrix:
[[83 16]
 [21 34]]
```

Calculate accuracy score,confusion matrix ,Create heatmap of the confusion matrix,# Print classification report

```
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
class_report = classification_report(y_test, y_pred)
print("\nClassification Report:")
print(class_report)
```

Accuracy: 0.7597

Confusion Matrix



Classification Report:
precision recall f1-score support